



# HUST-USYD Summer School on Parallel Programming Practice – Lecture 3

Bing Bing Zhou ([bing.zhou@sydney.edu.au](mailto:bing.zhou@sydney.edu.au))

School of Computer Science, University of Sydney

# Outline

- Review of Homework 1
- Speedup and efficiency
- Amdahl's law and overheads
- General design process
  - Machine independent part
    - Partitioning, Communication/synchronization
    - Task dependency graph
    - Math associative law
  - Machine dependent part
    - Assignment and load balancing
- Lab exercise: Gaussian elimination with partial pivoting

# Speedup and Efficiency

- Parallel computing is for high performance – high speed
- Speedup of a parallel algorithm is a measure of relative performance improvement over sequential algorithms for solving a given problem
  - Defined as the ratio of the compute time of a fastest sequential algorithm over the time of a parallel algorithm
- Let  $T_s$  be the compute time using a single processor and  $T_p$  be the time using  $p$  processors. The speedup is then defined as

$$S = \frac{T_s}{T_p}$$

- Efficiency of a parallel algorithm is the proportion of processors usefully utilized by the computation and defined as


$$E = \frac{S}{p} = \frac{T_s}{pT_p}$$

# Speedup and Efficiency

- When using  $p$  processors to solve a problem, do we expect  $p$  speedup, i.e., high efficiency?
  - Probably not!
  - Overheads in addition to the computation will be introduced in most parallel programs and they include
    - Process/thread communication or synchronization
    - Workload imbalance among available processors/threads
    - Extra work introduced to manage the computation and increase parallelism
    - ...
- Let  $T_o = pT_p - T_s$  be the total overhead. Then
$$E = \frac{T_s}{pT_p} = \frac{T_s}{T_o + T_s} = \frac{1}{1 + \frac{T_o}{T_s}}$$
- $0 \leq E \leq 1$  and  $E$  will be small when  $T_o$  is large


# Ideal Speedup

- Speedup:  $S = \frac{T_s}{T_p}$
- Ideally, we expect the speedup is increased by a factor of  $p$  when using  $p$  processors



# Actual Speedup

- In practice, the actual speedup for solving a problem is often smaller than the ideal speedup and worse as the number of processors increases



# Amdahl's Law

- Amdahl's law is used to predict the theoretical speedup when using multiple processors for parallel computing
- It shows that the serial parts of a parallel program impose a hard limit on potential speedup
- Assume the total amount of operations for solving a problem can be divided into two parts:
  - One part  $\beta$  is purely sequential
  - The other part  $1 - \beta$  is perfectly parallelizable
- The parallel time using  $p$  processors will be

$$T_p = \beta T_s + (1 - \beta)T_s/p$$

# Amdahl's Law

- The speedup is

$$S = \frac{T_s}{T_p} = \frac{T_s}{\beta T_s + (1 - \beta)T_s/p} = \frac{p}{1 + \beta(p - 1)}$$

- When  $p$  is very large, we have

$$S \rightarrow \frac{1}{\beta}$$

- Then  $\beta$  becomes a limiting factor

- E.g., when only 5% of the program are sequential , i.e.,  $\beta = 0.05$ , no matter how many processors are used, the speedup cannot be greater than 20!



# Overheads

- We know

$$E = \frac{T_s}{pT_p} = \frac{T_s}{T_o + T_s} = \frac{1}{1 + \frac{T_o}{T_s}}$$

- Main cause of Inefficiency in addition to poor single processor performance is overheads
- **Therefore, we must try the best to minimize all unnecessary overheads in designing and implementing parallel algorithms**

# Designing Parallel Algorithms

- How a problem specification is translated into an algorithm that displays concurrency, scalability, and locality
- Parallel algorithm design is not easily reduced to simple recipes
- Rather, it requires the sort of integrative thought that is commonly referred to as ``creativity“
- May need new ideas that have not been studied before

# General Design Process

- General design process involves two stages:
  - Machine independent stage
    - Recognize opportunities for parallel execution based on the characteristics of a given problem
      - Partitioning: divide a large task into multiple smaller ones which can be executed concurrently
      - Communication/synchronization: coordinate the execution of concurrent tasks and establish appropriate communication/synchronization structures
  - Machine dependent stage
    - Assignment: reorganize tasks and assign them to multiple processes/threads based on characteristics of a specific machine
      - minimize overheads and balance workload across processors

# Partitioning


- Expose opportunities for parallel execution
- The focus is on defining a large number of small tasks, each of which consists of the computation and the data on which this computation operates
- Typical types of partitioning:
  - Task partitioning
    - Divide the computation into pieces first
    - Then associate data with the computations
  - Data partitioning
    - Divide data into pieces first
    - Then associate computations with the data
- Which one should be applied depends on the actual problem

# Communication/Synchronization

- The tasks generated by a partition are intended to execute concurrently but cannot, in general, execute independently
  - Data must then be transferred between tasks so as to allow computation to proceed
- Communication/synchronization is then required to manage the data transfer and/or coordinate the execution of tasks
- Organizing communication in an efficient manner can be challenging


## Example: Matrix-vector multiplication

- Task partitioning:
  - Partition vector  $y$ , one element per task
  - Computing each element  $c$  involves one row of  $A$  and vector  $b$
- Observations:
  - Task size is uniform
  - No dependences between tasks
  - Embarrassingly parallel



# Example: Matrix-vector multiplication

- Data partitioning:
  - We can partition matrix A and vector b and then associate one multiplication with each pair of data items, one from A and one from b as a task
  - Note the results from these small tasks are just intermediate results, which will be considered as new data for further new task construction
  - This is equivalent to further partition each inner product into n smaller tasks
- Observations:
  - Task size is uniform
  - Good for fine-grained parallelism
  - But dependences between tasks for inner product of two vectors!
    - Require comm/sync



# Parallel Structure of Algorithm

- To design parallel algorithms for solving a given problem, typically the first step is to detect parallel structures of the sequential algorithms
- If the algorithm's parallel structure is determined, many subsequent decisions become obvious
- Task-dependency graph is one of the good techniques to identify program's parallel structures




# Task-Dependency Graph

- In a task-dependency graph each node represents a task and the arrows between nodes indicate dependencies between tasks
- The main purpose is to show the parallel structure to demonstrate various properties of the algorithm, e.g.,
  - parallel structure patterns – regularity
  - edge set features - dependency
  - structure dependent on input data – work prediction
- It is often enough to present one small graph
- However, it is challenging to construct the graph for complex algorithms

# Task-Dependency Graph

- Matrix-vector multiplication:
  - Each node represents a multiplication of  $a_{ij}$  and  $b_j$
  - Arrows indicate sequential additions (data dependency)
  - Each vertical arrow-node line represents an output of dot product  $y_i$



## Math Associative Law

- The parallel structure of an algorithm is an important concept, but it should not be used alone
- To evaluate the parallelism potential of an algorithm, the mathematics behind the algorithm plays an equally important role
- Knowing the mathematical basics of an algorithm can increase the degree of parallelism – very important in practice

## Math Associative Law

- Consider summation of  $n$  integers:  $s = \sum_{i=0}^{n-1} a_i$

- For a typical sequential program:

$s = 0;$

for ( $i=0; i<n; i++$ )

$s = s + a[i];$


- The structure of the algorithm:



- It is completely sequential in nature!


# Math Associative Law

- Addition operation obeys the associative law
- For the summation, additions can be done in any order
- We can then have a new algorithm:  
$$\sum_{i=0}^{2n-1} a_i = \sum_{i=0}^{n-1} a_i + \sum_{i=n}^{2n-1} a_i$$
  - Pairwise additions recursively
- The parallel structure forms a binary tree
  - Degree of parallelism changes from  $n$  to 1



# Math Associative Law

- Divide and conquer technique:
  - Divide: partition data into  $n$  groups, each having one element
  - Conquer: pairwise addition recursively
- Technique is widely used in practice for reduction operations
- The operator not necessarily addition, any operation which obeys associative law



# Assignment

- After the structure and degree of parallelism for a problem are identified, we can construct tasks based on the characteristics of specific machines and assign tasks to processors
- Question: What constitutes a task?
- Three levels of parallelism:
  - Instruction level
    - Fine granularity for ILP
  - Thread level
    - Coarse granularity for shared-memory machines to reduce overhead for synchronization
    - Fine granularity for GPU to make all cores busy
  - Process level
    - Coarse granularity partitioning for distributed-memory machines to reduce communication costs
- The main purpose of assignment is to
  - minimize communication /synchronization costs and
  - balance workload across the processes/threads

# Assignment

- The quality of task assignment is directly related to the quality of parallel computation
  - Some problems, e.g., dense matrix computation, regular meshes, can be easily decomposed into a number of tasks of equal size with regular data dependency pattern
    - task assignment relatively easy
  - Some other problems, e.g., sparse matrices, unstructured meshes, and graphs, are more complicated, during the computation we need to consider how to
    - balance workload
    - minimize communication/synchronization overhead
    - redundant computation



## Lab exercise: Gaussian Elimination with Unrolling

- Revise program `gepp_0.c` to add loop unrolling with unrolling factor = 4 for general cases, i.e.,  $n$  may not be divisible by 4
- Test the correctness of your program and check the performance
- If unable to complete, take it home as Homework 2

# Gaussian Elimination

- For a given matrix A of size N by N, add multiples of each row to later rows to make A upper triangular:

//for each column l zero it out below the diagonal by adding

//multiples of row i to later rows

for i = 1 to n-1

// for each row j below row i


for j = i+1 to n

// add a multiple of row i to row j


tmp = A(j,i);

for k = i to n


$$A(j,k) = A(j,k) - (\text{tmp}/A(i,i)) * A(i,k)$$



After i=1




After i=2



After i=3

...




After i=n-1

# Gaussian Elimination

- Store multipliers  $m$  below diagonal in zeroed entries for later use

```
for i = 1 to n-1
  for j = i+1 to n
     $A(j,i) = A(j,i)/A(i,i)$ 
    for k = i+1 to n
       $A(j,k) = A(j,k) - A(j,i) * A(i,k)$ 
```



Store  $m$  here

- Call the strictly lower triangular matrix of multipliers  $M$ , and let  $L = I+M$
- Call the upper triangle of the final matrix  $U$
- *Lemma (LU Factorization)*: If the above algorithm terminates (does not divide by zero) then  $A = L*U$

# Gaussian Elimination with Partial Pivoting

- $A = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$  fails completely because can't divide by  $A(1,1)=0$
- But solving  $Ax=b$  should be easy!
- When diagonal  $A(i,i)$  is tiny (not just zero), algorithm may terminate but get completely wrong answer
  - Numerical instability
  - Roundoff error is cause
- Cure: Pivot (swap rows of  $A$ ) so  $A(i,i)$  large

# Gaussian Elimination with Partial Pivoting

– Partial Pivoting: swap rows so that  $A(i,i)$  is largest in column for  $i = 1$  to  $n-1$

find and record  $k$  where  $|A(k,i)| = \max_{\{i \leq j \leq n\}} |A(j,i)|$

... i.e. largest entry in rest of column  $i$

if  $|A(k,i)| = 0$

exit with a warning that  $A$  is singular, or nearly so

elseif  $k \neq i$

swap rows  $i$  and  $k$  of  $A$

end if

$A(i+1:n,i) = A(i+1:n,i) / A(i,i)$  ... each  $|\text{quotient}| \leq 1$

$A(i+1:n,i+1:n) = A(i+1:n, i+1:n) - A(i+1:n, i) * A(i, i+1:n)$

