



HUST-USYD Summer School on Parallel Programming Practice – Lecture 5

Bing Bing Zhou (bing.zhou@sydney.edu.au)

School of Computer Science, University of Sydney

Outline

- Synchronization
 - Race condition and mutual exclusion
- OpenMP synchronization constructs
 - Critical, atomic and barrier
- Data dependency
 - Instruction-level data dependency
 - Loop-carried data dependency
- Reduction operation
- OpenMP reduction clause
- Exercise: reduction operation
- Scan (or prefix) operation
- Homework 3

Synchronization

- In shared-memory machines global data are shared by multiple threads
- Synchronization must be used to
 - Protect access to shared data and so to prevent race conditions
 - Impose order constraints
- OpenMP supported synchronization constructs
 - **Critical**, **atomic** and **barrier**
 - The full OpenMP specification has several more

Synchronization Constructs

- OpenMP **critical** directive:
- **#pragma omp critical**
- Threads wait their turn – only one at a time can execute in the **critical** construct
- Critical directive must be inside a **parallel** region
- E.g., threads to increment count

```
#pragma omp parallel shared(count)  
{  
    ...  
    #pragma omp critical  
    count++;  
    ...  
}
```

Synchronization Constructs

- OpenMP **atomic** directive:
- **#pragma omp atomic**
- Also provides mutual exclusion, but only applies to the update of one memory location
- **atomic** directive applies only to the statement immediately following it
- E.g., accumulate partial results done by multiple threads

```
double sum = 0.0;
```

```
#pragma omp parallel shared(sum)
```

```
{
```

```
    double lsum;
```

```
    ...
```

```
#pragma omp atomic
```

```
sum += lsum
```

```
    ...
```

```
}
```

Synchronization Constructs

- OpenMP **barrier** directive:
- **#pragma omp barrier**
- **barrier** is a point in a program all threads must reach before any threads are allowed to proceed
- In OpenMP most constructs have an implied barrier at the end of the constructs
 - To remove that barrier, use **onwait** clause

```
#pragma omp parallel shared(sum)
```

```
{
```

```
    //all threads do something
```

```
    #pragma omp barrier //Threads wait until all threads hit the barrier
```

```
    //all threads do some other thing
```

```
}
```

Instruction-Level Data Dependency

- Data dependency at the instruction level – the dependency in a sequence of instructions:
 - Flow dependency (Read After Write or RAW)
 - E.g., $a = b + c$ (Write $b+c$ to a)
 $d = a * e$ (Read a out and assign $a*e$ to d)
 - Anti-dependency (WAR)
 - E.g., $b = a + c$ (Read a out and assign $a+c$ to b)
 $a = d * e$ (Write $d+e$ to a)
 - Output dependency (WAW)
 - E.g., $a = b + c$
 $a = d * e$

Instruction-Level Data Dependency

- Data dependency at the instruction level – the dependency in a sequence of instructions:
 - Flow dependency is true dependency
 - Anti-dependency and output dependency are not real (so called artificial) dependency and can be eliminated
 - E.g., output dependency in the previous example can be removed by writing d^*e to another variable and then use the new variable (instead of a) in the future

Loop-Carried Data Dependency

- Can we parallelize the following for loop?

```
for (i=0; i<n; i++)
```

```
{
```

```
    tmp = a[i];
```

```
    a[i] = b[i];
```

```
    b[i] = tmp;
```

```
}
```

- Clearly there are data dependencies in the three instructions
- However, there are no data dependencies to be carried out to the next iteration
- Therefore, the loop can be safely parallelized
- In OpenMP we really need to check loop-carried dependency!

Loop-Carried Data Dependency

- Can we parallelize the following for loop?

```
for (i=0; i<n; i++)
```

```
{
```

```
    a[i] = d[i] + e[i];
```

```
    d[i] = e * a[i+1];
```

```
}
```

- Not very clear?
- Let's expand the for loop starting from i=0

- a[i] in red means assigned **new** value
- a[i] in green means **old** value

WAR
dependency

a[0] = d[0] + e[0]

d[0] = e * a[1]

a[1] = d[1] + e[1]

d[1] = e * a[2]

a[2] = d[2] + e[2]

d[2] = 3 * a[3]

a[3] = d[3] + e[3]

d[3] = e * a[4]

⋮

Loop-Carried Data Dependency

- Now consider parallelize the for loop:

Thread 0:

$$a[0] = d[0] + e[0]$$

$$d[0] = e * a[1]$$

$$a[1] = d[1] + e[1]$$

$$d[1] = e * a[2]$$

Thread 1:

$$a[2] = d[2] + e[2]$$

$$d[2] = 3 * a[3]$$

$$a[3] = d[3] + e[3]$$

$$d[3] = e * a[4]$$

Thread 2:

$$a[4] = d[0] + e[0]$$

$$d[0] = e * a[5]$$

$$a[5] = d[1] + e[1]$$

$$d[1] = e * a[6]$$



Wrong values!

- Now it becomes clear – there are indeed loop-carried data dependencies
- Also potential data race
 - i.e., one thread write and another read on the same data
- How to deal with this problem?

Loop-Carried Data Dependency

```
for (i=0; i<n; i++)  
{
```

```
    a[i] = d[i] + e[i];
```

```
    d[i] = e * a[i+1];
```

```
}
```



```
for (i=0; i<n; i++)  
{
```

```
    d[i] = e * a[i+1];
```

```
    a[i] = d[i] + e[i];
```

```
}
```



```
for (i=0; i<n; i++)
```

```
    d[i] = e * a[i+1];
```

```
for (i=0; i<n; i++)
```

```
    a[i] = d[i] + e[i];
```

- In the same i iteration, the order of two instructions can be exchanged because there is no data dependency
- Also separate the instructions into two for loops

Loop-Carried Data Dependency

- Now parallelize two for loops one by one:

```
#pragma omp parallel for
```

```
for (i=0; i<n; i++)
```

```
    d[i] = e * a[i+1];
```

Implied barrier here



```
#pragma omp parallel for
```

```
for (i=0; i<n; i++)
```

```
    a[i] = d[i] + e[i];
```

Thread 0:

$d[0] = e * a[1]$

$d[1] = e * a[2]$

$a[0] = d[0] + e[0]$

$a[1] = d[1] + e[1]$

Thread 1:

$d[2] = 3 * a[3]$

$d[3] = e * a[4]$

$a[2] = d[2] + e[2]$

$a[3] = d[3] + e[3]$

Thread 3:

$d[4] = e * a[5]$

$d[5] = e * a[6]$

$a[4] = d[0] + e[0]$

$a[5] = d[1] + e[1]$

- We now obtain the correct results!

Loop-Carried Data Dependency

- Can we parallelize the following for loop?

```
for (i=0; i<n; i++)
```

```
{
```

```
    a[i+1] = d[i] + e[i];
```

```
    d[i] = e * a[i];
```

```
}
```

- A simple problem for you to work out at home

Loop-Carried Data Dependency

- In the previous lecture we also discussed how to deal with another type of loop-carried data dependency:

```
int i, j, A[MAX];
j = 5;
for (i=0; i< MAX; i++){
    j +=2;
    A[i] = big(j);
}
```

Note: loop index “i” is private by default

Remove loop carried dependence

```
int i, A[MAX];
#pragma omp parallel for
for (i=0; i< MAX; i++) {
    int j = 5 + 2*(i+1);
    A[i] = big(j);
}
```

- In general we need to pay attention on the following loop-carried data dependency:
 - index of an array data is not equal to current loop index
 - The value of a variable **changes with the iterations**

Reduction Operation

- Given associative operator \oplus
- Examples
 - Add (+)
 - Multiply (*)
 - And, Or (&&, ||)
 - max, min
- and also an array of elements $[a_0, a_1, a_2, \dots, a_{n-1}]$
- Calculate
- $S = a_0 \oplus a_1 \oplus a_2 \oplus \dots \oplus a_{n-1}$

Reduction Operation

- E.g., let $\oplus = +$, that is, summation of n numbers:
 - $S = a_0 + a_1 + a_2 + \dots + a_{n-1} = \sum_{i=0}^{n-1} a_i$
- Use recursion:
 - $S_0 = a_0$
 - $S_i = S_{i-1} + a_i$ for $1 \leq i < n$
- A sequential routine:
$$S = a[0];$$
$$\text{for } (i=1; i < n; i++)$$
$$S = S + a[i];$$
- Can we simply use OpenMP for directive to parallelize this sequential routine?
 - No, it is recursive and the value of S changes with iterations
 - There are loop-carried data dependencies!

Reduction Operation

- E.g., let $\oplus = +$, that is, summation of n numbers:
 - $S = a_0 + a_1 + a_2 + \dots + a_{n-1} = \sum_{i=0}^{n-1} a_i$
- However, addition is associative and commutative
 - We can add the numbers in any order
 - E.g., $S = \sum_{i=0}^{k-1} a_i + \sum_{i=k}^{2k-1} a_i + \sum_{i=2k}^{n-1} a_i$
 $= S_0 + S_1 + S_2$
- Then we can use the OpenMP directive to parallelize the for loop for each thread to compute a partial sum and then accumulate them to obtain the final result

Reduction Operation

- E.g., let $\oplus = +$, that is, summation of n numbers:
 - $S = a_0 + a_1 + a_2 + \dots + a_{n-1} = \sum_{i=0}^{n-1} a_i$

```
#pragma omp parallel shared(a,S) private(i)
```

```
{
```

```
    double lsum = 0.0; //each thread has a local lsum variable
```

```
    #pragma omp for nowait
```

```
    for (i=0; i<n; i++) //the number of iterations distributed
```

```
        lsum += a[i]; //each thread calculate a partial sum
```

```
    ... //accumulate partial sums
```

```
}
```

- How to accumulate the partial sums?

OpenMP Reduction Clause

- OpenMP supports a reduction clause:
 - **reduction (op : list)**
- E.g.,
double S=0;
#pragma omp parallel for shared(S, a), reduction(+:S)
for (int i=0; i<n; ++)
 S += a[i];
- When reduction clause is added
 - A local copy of each list variable is made and initialized, depending on the “op” (e.g. 0 for “+”)
 - Updates occur on the local copy
 - Local copies are then reduced into a single value
 - The variables in “list” must be shared in the enclosing parallel region

OpenMP Reduction Clause

- Many different associative operands can be used
- Initial values are the ones that make sense mathematically

Operator	Initial value	Operator	Initial value
+	0	&&	1
*	1		0
-	0	&	~0
min	Largest pos. number		0
max	Most neg. number	^	0

Lab Exercise: Reduction Operation

- Write an OpenMP program to compute summation of n numbers
- You need to write two routines:
 - One uses OpenMP reduction clause
 - The other is to let each thread compute a partial sum and then accumulate the partial sums
- Run your program to compare the performance of these two routines

Scan Operation

- Let $A = [a_0, a_1, a_2, \dots, a_{n-1}]$ be an array of elements and \oplus an associative operator
- Scan (or prefix) operation produces another array C :
$$C = [a_0, (a_0 \oplus a_1), (a_0 \oplus a_1 \oplus a_2), \dots, (a_0 \oplus a_1 \oplus a_2 \oplus \dots \oplus a_{n-1})]$$
- E.g., $A = [1, 2, 3, 4, 5, 6]$, and after add scan of A we obtain
 $C = [1, 3, 6, 10, 15, 21]$
- It can be seen that there are many redundancy in calculating C_i
 - Use recursion to remove redundant operations
- A sequential routine:
$$c[0] = a[0];$$
$$\text{for } (i=1; i < n; i++)$$
$$c[i] = c[i-1] \oplus a[i];$$
- It takes only $n - 1 \oplus$ operations

Parallel Algorithm for Scan Operation

- The sequential routine is very simple, but strictly sequential in nature
 - True data dependency and cannot directly use OpenMP for directive to parallelize the for loop
- Need to design new parallel algorithms
- Many parallel algorithms for scan operation have been developed based on the fact that \oplus is an associative operator
 - Key is to minimize the amount of operations
- In the following we discuss one parallel algorithm which is efficient for parallel computation of scan operation on shared-memory machines

Parallel Algorithm for Scan Operation

– Rewrite the c_i :

$$c_0 = a_0$$

$$c_1 = a_0 \oplus a_1$$

$$c_2 = a_0 \oplus a_1 \oplus a_2$$

$$c_3 = a_0 \oplus a_1 \oplus a_2 \oplus a_3$$

$$c_4 = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4$$

$$c_5 = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5$$

$$c_6 = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5 \oplus a_6$$

$$c_7 = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5 \oplus a_6 \oplus a_7$$

$$c_8 = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5 \oplus a_6 \oplus a_7 \oplus a_8$$

$$c_9 = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5 \oplus a_6 \oplus a_7 \oplus a_8 \oplus a_9$$

$$c_{10} = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5 \oplus a_6 \oplus a_7 \oplus a_8 \oplus a_9 \oplus a_{10}$$

$$c_{11} = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5 \oplus a_6 \oplus a_7 \oplus a_8 \oplus a_9 \oplus a_{10} \oplus a_{11}$$

Parallel Algorithm for Scan Operation

- Reduce redundant calculation:

$$c_0 = a_0$$

$$c_1 = a_0 \oplus a_1$$

$$c_2 = a_0 \oplus a_1 \oplus a_2$$

$$c_3 = a_0 \oplus a_1 \oplus a_2 \oplus a_3$$

$$c_4 = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4$$

$$c_5 = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5$$

$$c_6 = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5 \oplus a_6$$

$$c_7 = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5 \oplus a_6 \oplus a_7$$

$$c_8 = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5 \oplus a_6 \oplus a_7 \oplus a_8$$

$$c_9 = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5 \oplus a_6 \oplus a_7 \oplus a_8 \oplus a_9$$

$$c_{10} = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5 \oplus a_6 \oplus a_7 \oplus a_8 \oplus a_9 \oplus a_{10}$$

$$c_{11} = a_0 \oplus a_1 \oplus a_2 \oplus a_3 \oplus a_4 \oplus a_5 \oplus a_6 \oplus a_7 \oplus a_8 \oplus a_9 \oplus a_{10} \oplus a_{11}$$

Divide the computation
into blocks

Parallel Algorithm for Scan Operation

– Reduce redundant calculation:

$$c_0 = a_0$$

$$c_1 = a_0 \oplus a_1$$

$$c_2 = a_0 \oplus a_1 \oplus a_2$$

$$c_3 = c_{0-2} \oplus a_3$$

$$c_4 = c_{0-2} \oplus a_3 \oplus a_4$$

$$c_5 = c_{0-2} \oplus a_3 \oplus a_4 \oplus a_5$$

$$c_6 = c_{0-2} \oplus c_{3-5} \oplus a_6$$

$$c_7 = c_{0-2} \oplus c_{3-5} \oplus a_6 \oplus a_7$$

$$c_8 = c_{0-2} \oplus c_{3-5} \oplus a_6 \oplus a_7 \oplus a_8$$

$$c_9 = c_{0-2} \oplus c_{3-5} \oplus c_{6-8} \oplus a_9$$

$$c_{10} = c_{0-2} \oplus c_{3-5} \oplus c_{6-8} \oplus a_9 \oplus a_{10}$$

$$c_{11} = c_{0-2} \oplus c_{3-5} \oplus c_{6-8} \oplus a_9 \oplus a_{10} \oplus a_{11}$$

$$c_{i-j} = a_i \oplus a_{i+1} \cdots \oplus a_j$$


$$\text{e.g., } c_{3-5} = a_3 \oplus a_4 \oplus a_5$$

$$c_{i-k} \oplus c_{(k+1)-j} = c_{i-j}$$

$$\text{e.g., } c_{0-2} \oplus c_{3-5} = c_{0-5}$$

Parallel Algorithm for Scan Operation

- Parallel algorithm:
- Stage 1: Each thread gets a subset of n/t elements (block partitioning), performs local scan operation on assigned elements, and then thread i (except the last one) stores its last local c element to a working array at $w[i + 1]$



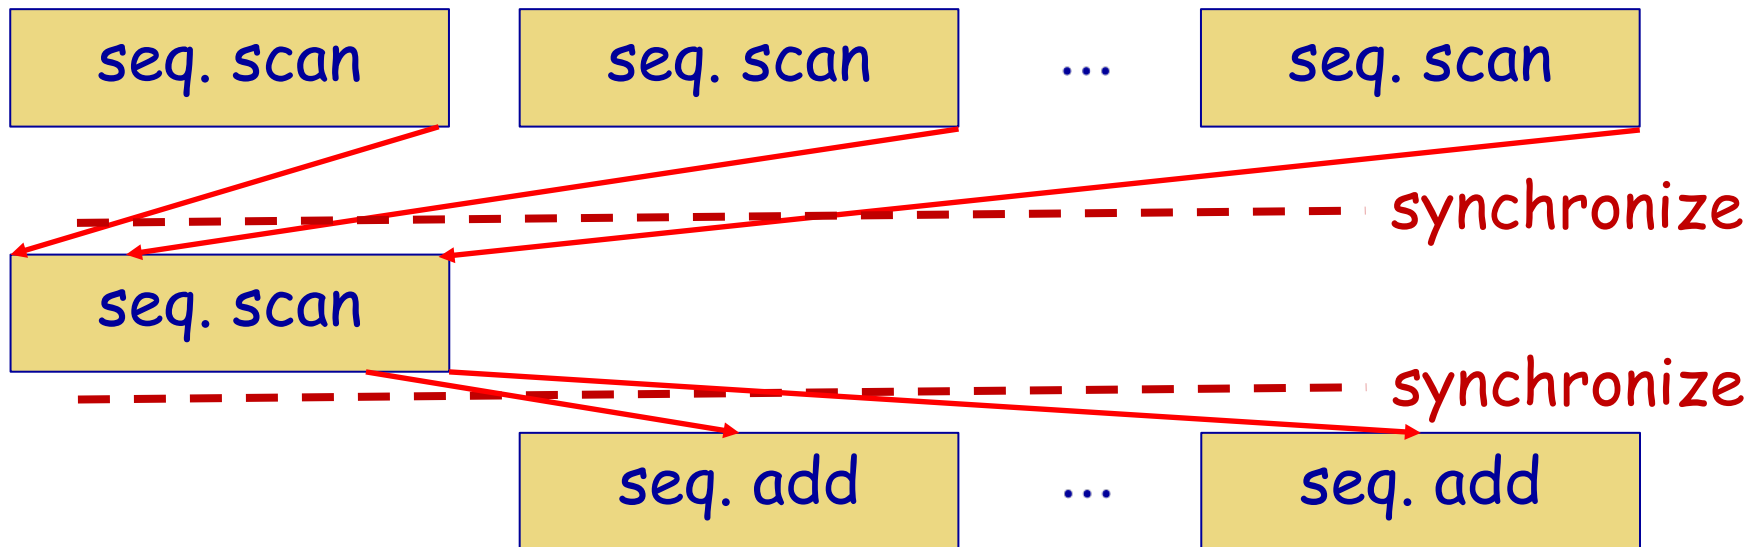
Sync must be applied (implicitly or explicitly) btw stages!

- Stage 2: A single thread performs scan operation on W
- Stage 3: Each thread $\oplus w[i]$ to every partial scan element

$c_0 = c_0 \oplus 0$		$c_3 = c_3 \oplus c_{0-2}$		$c_6 = c_6 \oplus c_{0-5}$
$c_1 = c_{0-1} \oplus 0$		$c_4 = c_{3-4} \oplus c_{0-2}$		$c_7 = c_{6-7} \oplus c_{0-5}$
$c_2 = c_{0-2} \oplus 0$		$c_5 = c_{3-5} \oplus c_{0-2}$		$c_8 = c_{6-8} \oplus c_{0-5}$

Parallel Algorithm for Scan Operation

- To summarize, this parallel algorithm consists of three parallel stages
 - Synchronization between stages
 - In 2nd stage only one thread is active
 - Purely sequential, but overhead not heavy as SM machines are usually small



Parallel Algorithm for Scan Operation

- Total amount of operations:
 - Stage 1: each thread performs local scan on n/t elements
 - $(n/t - 1) * t = n - t$
 - Stage 2: a single thread performs scan on W
 - $t - 1$ (the size of W is t)
 - Stage 3: each thread add w_i to the partially scanned elements
 - $n/t * t = n$
- Thus the total amount is
 - $n - t + t - 1 + n = 2n - 1$
- The total amount of operations in sequential computation is $n - 1$
- Thus the parallel algorithm takes less amount of time to complete the computation if the number of threads > 2

Homework 3: Parallel Scan Operation

- Implement the parallel algorithm for scan operation using OpenMP
 - Compare the performance with the sequential algorithm
 - In a **parallel** region, if we want part of code to be done by a single (any) thread, we can use **single** directive
- #pragma omp single**
- All other threads will skip the **single** region and stop at the **barrier** at the end of the **single** construct until all threads have reached the barrier

