



HUST-USYD Summer School on Parallel Programming Practice – Lecture 9

Bing Bing Zhou (bing.zhou@sydney.edu.au)

School of Computer Science, University of Sydney

Outline

- Non-blocking send/recv
- Collective communication routines
- Lab exercise: Matrix-vector multiplication revisited
- Parallel matrix multiplication on 2D mesh/torus
- MPI Cartesian topology routines
- Homework 7

Non-Blocking Send & Receive

- MPI Functions: `MPI_Isend`, `MPI_Irecv`
- `int MPI_Isend(const void* buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm, MPI_Request *request)`
- `int MPI_Irecv(void* buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Request *request)`
- When `MPI_Isend/MPI_Irecv` is called, the calling process will be not blocked and it continues to execute the subsequent instructions
- `MPI_Request *request`: a so called opaque object, which identifies communication operations and matches the operation that initiates the communication with the operation that terminates it

Non-Blocking Send & Receive

- MPI Functions: MPI_Test, MPI_Wait
- `int MPI_Test(MPI_Request *request, int *flag, MPI_Status *status)`
- `int MPI_Wait(MPI_Request *request, MPI_Status *status)`
- MPI_Test tests if operation finished
 - If not, the calling process will continue its execution
- MPI_Wait blocks the calling process until the non-blocking send/rcv operation is finished

Non-Blocking Send & Receive


- MPI Functions: `MPI_Testsome`, `MPI_Waitsome`
- `int MPI_Testsome(int incount, MPI_Request *array_of_requests, int *outcount, int *array_of_indices, MPI_Status *array_of_statuses)`
- `int MPI_Waitsome(int incount, MPI_Request *array_of_requests, int *outcount, int *array_of_indices, MPI_Status *array_of_statuses)`
- Tests or waits until at least one of the operations associated with active handles in the list have completed
- Needs an array of requests (also indices and statuses) for multiple `Isend/Irecv` operations

Collective Communication

- MPI specifies a set of collective communication functions
 - One-to-All: One process contributes to the result and all processes receive the result
 - `MPI_Bcast`, `MPI_Scatter`, `MPI_Scatterv`, ...
 - All-to-One: All processes contribute to the result and one process receives the result
 - `MPI_Gather`, `MPI_Gatherv`, ...
 - All-to-All: All processes contribute to the result and all processes receive the result
 - `MPI_Allgather`, `MPI_Alltoall`, `MPI_Allgatherv`, ...
 - Other: `MPI_Barrier`, `MPI_Reduce`, `MPI_Scan`, ...

Broadcast


- The MPI broadcast routine (one-to-all) is:
`int MPI_Bcast(void *buf, int count, MPI_Datatype datatype,
int source, MPI_Comm comm)`
- All processes must call `MPI_Bcast()`



Gather

- The MPI gather operation (all-to-one) is:

```
int MPI_Gather(void *sendbuf, int sendcount,  
              MPI_Datatype senddatatype, void *recvbuf,  
              int recvcount, MPI_Datatype recvdatatype,  
              int target, MPI_Comm comm)
```



Gather

- The MPI gather operation (all-to-one) is:


```
int MPI_Gather(void *sendbuf, int sendcount,  
              MPI_Datatype senddatatype, void *recvbuf,  
              int recvcount, MPI_Datatype recvdatatype,  
              int target, MPI_Comm comm)
```

- **sendcount** – number of elements each process sends
- **recvcount** – number of elements the target process will receive from each process
- ***recvbuf** – the length of the recvbuf must be at least recvcount times the number of processes involved in the operation

Allgather

- The MPI All gather operation (all-to-all) is:


```
int MPI_Allgather(void *sendbuf, int sendcount,  
                 MPI_Datatype senddatatype, void *recvbuf,  
                 int recvcount, MPI_Datatype recvdatatype,  
                 MPI_Comm comm)
```



Gatherv

- The MPI gatherv operation (all-to-one) is:

```
int MPI_Gatherv(void *sendbuf, int sendcount,  
               MPI_Datatype sendtype, void *recvbuf,  
               int *recvcounts, int *displs, MPI_Datatype recvttype,  
               int root, MPI_Comm comm)
```



Gatherv

- The MPI gatherv operation (all-to-one) is:


```
int MPI_Gatherv(void *sendbuf, int sendcount,  
                MPI_Datatype sendtype, void *recvbuf,  
                int *recvcounts, int *displs, MPI_Datatype recvtype,  
                int root, MPI_Comm comm)
```

```
int sendarray[100];  
int stride = 150;  
int *displs, *rcounts;  
...  
rbuf = (int *) malloc(numprocs*stride*sizeof(int));  
for (int i=0; i<numprocs; i++){  
    displs[i] = i * stride;  
    rcounts[i] = 100;  
}  
MPI_Gatherv(sendarray, 100, MPI_INT, rbuf, rcounts, displs, MPI_INT, root,  
            MPI_COMM_WORLD);
```

Scatter

- The MPI scatter operation (one-to-all) is:

```
int MPI_Scatter(void *sendbuf, int sendcount,  
               MPI_Datatype senddatatype, void *recvbuf,  
               int recvcount, MPI_Datatype recvdatatype,  
               int source, MPI_Comm comm)
```



Scatter

- The MPI scatter operation (one-to-all) is:

```
int MPI_Scatter(void *sendbuf, int sendcount,  
               MPI_Datatype senddatatype, void *recvbuf,  
               int recvcount, MPI_Datatype recvdatatype,  
               int source, MPI_Comm comm)
```

- **sendcount** – number of elements the source process will send to each process
- ***sendbuf** – the length of the sendbuf must be at least sendcount times the number of processes involved in the operation
- **recvcount** – number of elements each process will receive

Scatterv and Alltoall

- The MPI scatterv operation (one-to-all) is:

```
int MPI_Scatterv(void *sendbuf, int *sendcounts, int *displs,  
                MPI_Datatype sendtype, void *recvbuf,  
                int recvcount, MPI_Datatype recvtype, int root,  
                MPI_Comm comm)
```

- The MPI all-to-all communication operation is:

```
int MPI_Alltoall(void *sendbuf, int sendcount,  
                MPI_Datatype senddatatype, void *recvbuf,  
                int recvcount, MPI_Datatype recvdatatype,  
                MPI_Comm comm)
```

Other Collectives

- The barrier synchronization operation:

```
int MPI_Barrier(MPI_Comm comm)
```

- The MPI reduction operation (all-to-one) is:

```
int MPI_Reduce(void *sendbuf, void *recvbuf, int count,  
               MPI_Datatype datatype, MPI_Op op, int target,  
               MPI_Comm comm)
```

- If the result of the reduction operation is needed by all processes, MPI provides:

```
int MPI_Allreduce(void *sendbuf, void *recvbuf, int count,  
                  MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```


Predefined Reduction Operations

Operation	Meaning	Datatypes
MPI_MAX	Maximum	C integers and floating point
MPI_MIN	Minimum	C integers and floating point
MPI_SUM	Sum	C integers and floating point
MPI_PROD	Product	C integers and floating point
MPI LAND	Logical AND	C integers
MPI_BAND	Bit-wise AND	C integers and byte
MPI_LOR	Logical OR	C integers
MPI BOR	Bit-wise OR	C integers and byte
MPI_LXOR	Logical XOR	C integers
MPI_BXOR	Bit-wise XOR	C integers and byte
MPI_MAXLOC	Max, value-location	Data-pairs
MPI_MINLOC	Min, value-location	Data-pairs

Reduction Operations

- The operation `MPI_MAXLOC` combines pairs of values (v_i, l_i) and returns the pair (v, l) such that v is the maximum among all v_i 's and l is the corresponding l_i (if there are more than one, it is the smallest among all these l_i 's)
- `MPI_MINLOC` does the same, except for minimum value of v_i

Value	15	17	11	12	17	11
Process	0	1	2	3	4	5

`MinLoc(Value, Process) = (11, 2)`

`MaxLoc(Value, Process) = (17, 1)`

Reduction Operations

- MPI datatypes for data-pairs used with the `MPI_MAXLOC` and `MPI_MINLOC` reduction operations

MPI Datatype	C Datatype
<code>MPI_2INT</code>	pair of ints
<code>MPI_SHORT_INT</code>	short and int
<code>MPI_LONG_INT</code>	long and int
<code>MPI_LONG_DOUBLE_INT</code>	long double and int
<code>MPI_FLOAT_INT</code>	float and int
<code>MPI_DOUBLE_INT</code>	double and int

Reduction Operations

– Example:

```
struct {  
    double val;  
    int    rank;  
} in, out;
```

```
in.val = ain;
```

```
in.rank = myrank;
```

```
MPI_Reduce( in, out, 1, MPI_DOUBLE_INT, MPI_MAXLOC, root,  
comm );
```

```
if (myrank == root) {
```

```
    aout = out.val;
```

```
    ind = out.rank;
```


```
}
```

Lab Exercise: Matrix-Vector Multiplication Revisited

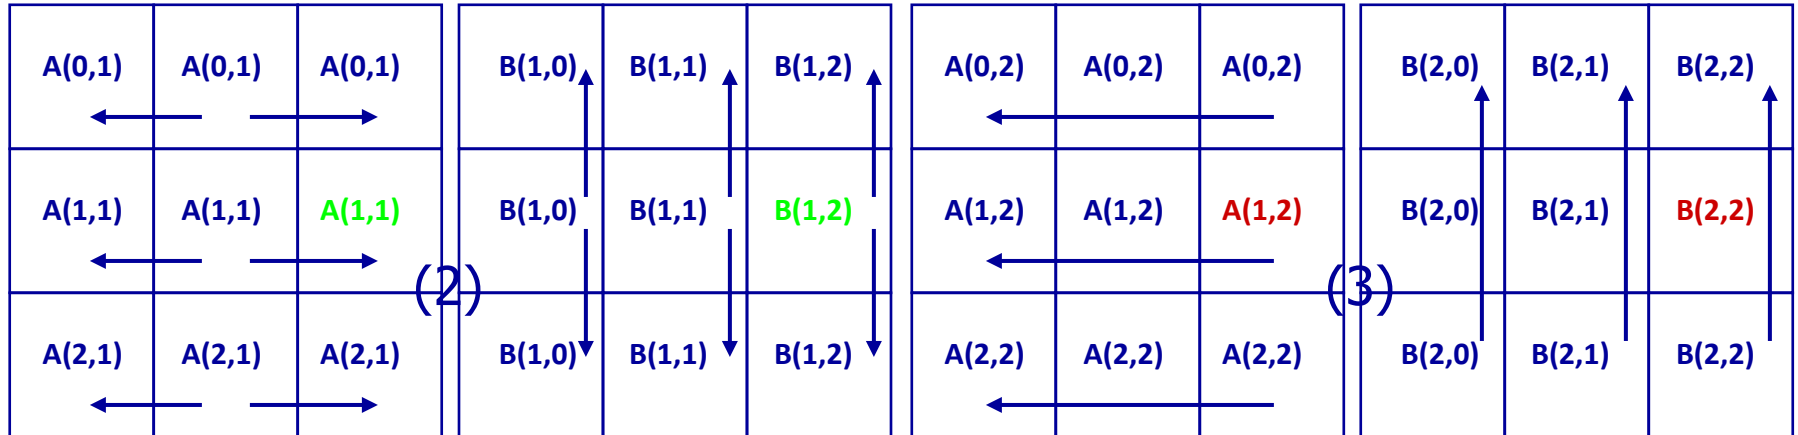
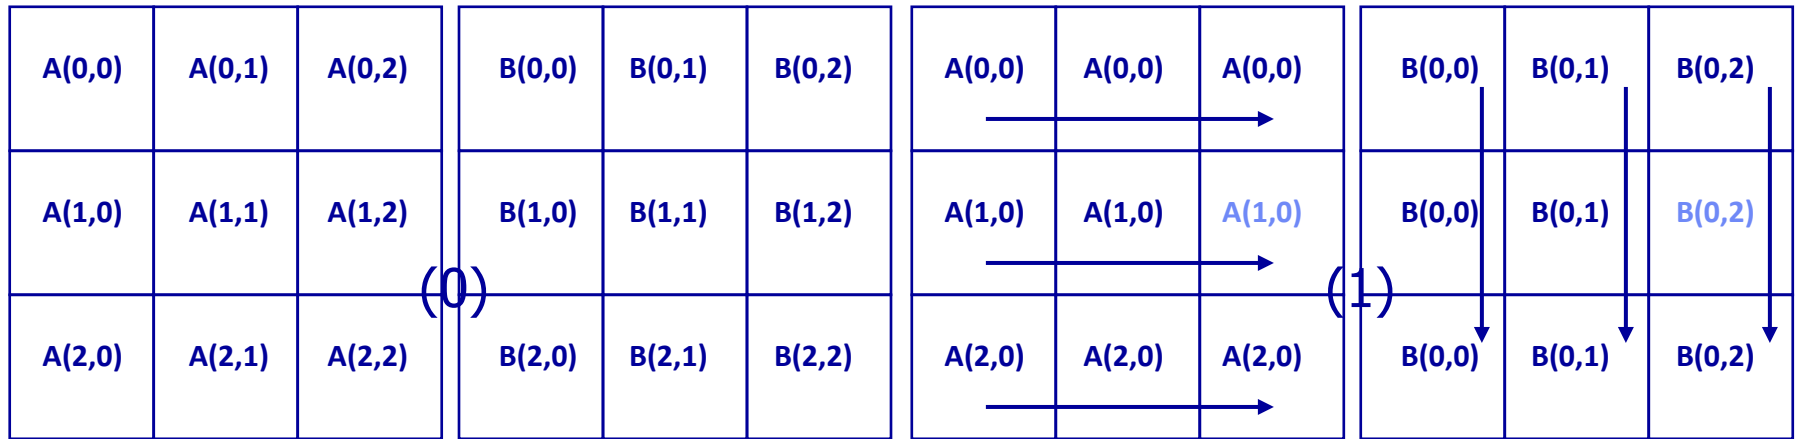
- Modify the program for matrix-vector multiplication you did in the last lecture by using collective communication routines
- You need to consider general cases, that is, n may not be divisible by p
- Thus you need to use `MPI_Scatterv` and `MPI_Gatherv` functions

Parallel MM on 2D Mesh/Torus

- Use 2D partitioning and assume $p = \sqrt{p} \times \sqrt{p}$, i.e., p processes are organized as a 2D mesh/torus
- Let $C(i, j)$ refer to a submatrix of size $n/\sqrt{p} \times n/\sqrt{p}$ and similar for $A(i, j)$ and $B(i, j)$
- Each $C(i, j)$ needs one row of $A(i, j)$ s and one column of $B(i, j)$ s which are held by different processes
- Communication: move $A(i, j)$ s horizontally and $B(i, j)$ s vertically

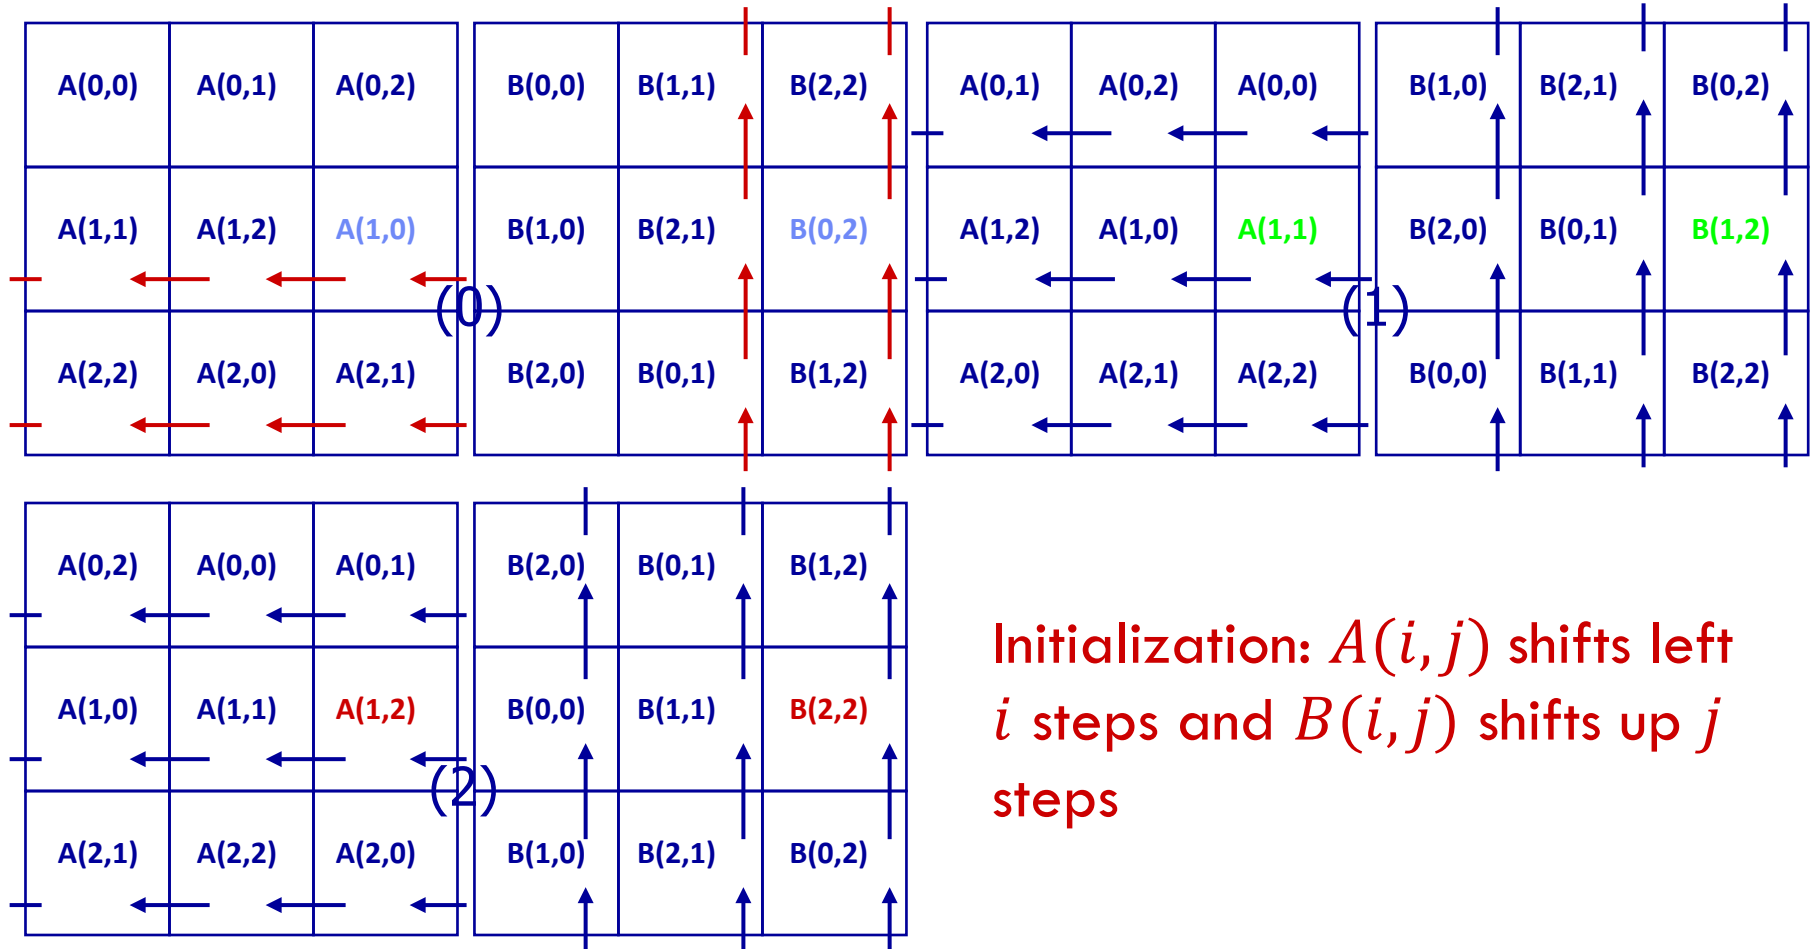


2D Algorithm 1 (Broadcast)



$$C(1,2) = A(1,0) * B(0,2) + A(1,1) * B(1,2) + A(1,2) * B(2,2)$$

2D Algorithm 1 (Cannon)



Initialization: $A(i, j)$ shifts left i steps and $B(i, j)$ shifts up j steps

$$C(1,2) = A(1,0) * B(0,2) + A(1,1) * B(1,2) + A(1,2) * B(2,2)$$

Parallel Matrix Multiplication

- We discussed 1D and 2D parallel algorithms for distributed-memory machines
- Question: which one performs better?
- 1D algorithm:
 - Matrices are partitioned into block rows ($n/p \times n$)
 - In one parallel step each process
 - Send one block row ($B(j)$) and received one block row ($\sim n^2/p$)
 - Multiplication of $A(i, j)$ ($n/p \times n/p$) and $B(j)$ ($n/p \times n$)
 - Number of operations is $\sim n^3/p^2$
 - Computational intensity: $\sim n/p$
 - Number of parallel steps is p

Parallel Matrix Multiplication

- 2D algorithm:
 - Matrices are partitioned into a number of submatrices $(n/\sqrt{p} \times n/\sqrt{p})$
 - In one parallel step each process
 - Send one submatrix $A(i, j)$ and receive one submatrix horizontally $(\sim n^2/p)$
 - Send one submatrix $B(i, j)$ and receive one submatrix vertically $(\sim n^2/p)$
 - Multiplication of $A(i, k)$ and $B(k, j)$
 - Number of operations is $\sim n^3/(p\sqrt{p})$
 - Computational intensity: $\sim n/\sqrt{p}$
 - Number of parallel steps is \sqrt{p}

Parallel Matrix Multiplication

- Although 2D algorithm requires to move both $A(i, j)$ and $B(i, j)$ in a parallel step, the number of operations increased by a factor of \sqrt{p}
- The computational intensity is increased by a factor of \sqrt{p}
- Since the total amount of work is fixed for a given problem, the number of parallel steps is also reduced by a factor of \sqrt{p}
 - Greatly reduced the communication overheads
- Therefore, 2D parallel algorithm for matrix multiplication

MPI Cartesian Topology Routines

- MPI allows a programmer to organize processors into logical k-D meshes
- The processor ids in `MPI_COMM_WORLD` can be mapped to other communicators (corresponding to higher-dimensional meshes) in many ways
- The goodness of any such mapping is determined by the interaction pattern of the underlying program and the topology of the machine

Cartesian Topologies

- To create Cartesian topologies using the function:

```
int MPI_Cart_create(MPI_Comm comm_old, int ndims, int *dims,  
int *periods, int reorder, MPI_Comm *comm_cart)
```

- This function takes the processes in the old communicator and creates a new communicator with ndims dimensions
- Each processor can now be identified in this new cartesian topology by a vector of dimension dims

Cartesian Topologies

- To determine process coords in cartesian topology given rank in group:

```
int MPI_Cart_coord(MPI_Comm comm_cart, int rank, int maxdims,  
                  int *coords)
```

- To determine process rank in communicator by its cartesian location:

```
int MPI_Cart_rank(MPI_Comm comm_cart, int *coords, int *rank)
```

Cartesian Topologies

- To find the resulting source and destination ranks, given a shift direction and amount

```
int MPI_Cart_shift( MPI_Comm comm, int direction, int displ,  
                  int *source, int *dest)
```

- To partition a cartesian topology into subgrids

```
int MPI_Cart_sub(MPI_Comm comm, const int remain_dims[],  
                MPI_Comm* new_comm)
```

- **MPI_Cart_sub** is a collective comm routine
 - All involved processes must call

Homework 7: MPI Cartesian Topology

- Search the Internet for detailed descriptions on MPI Cartesian topology
- Then write a simple MPI program to create 16 processes
- Use MPI cartesian routine to organize the processes into a 4 X 4 2D ring
- Declare two integers a and b
 - Assign process's id to a and then every process shifts a to the left neighbour process using MPI send/recv
 - Assign process's id to b and then every process shifts b up to the neighbor process above also using MPI send/recv
 - Print and check the results
- Partition processes into row and column subgrids
 - For processes with column id = 0, assign it original id to a and broadcast it to the processes in the same row
 - For processes with row id = 0, assign it original id to b and broadcast it to the processes in the same column
 - Print and check the results

